

# Oversampling and Undersampling Techniques

## Explaining techniques to handle imbalanced Dataset

In [29]:

```
import os
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import warnings
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import make_pipeline
from pylab import rcParams
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.metrics import precision_score, recall_score, confusion_matrix
from sklearn.metrics import f1_score, roc_auc_score, roc_curve
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
os.chdir('/Users/subham/Desktop/SMOTE')
```

In [30]:

```
%matplotlib inline
np.random.seed(27)
rcParams['figure.figsize'] = 10, 6
warnings.filterwarnings('ignore')
sns.set(style="darkgrid")
```

In [31]:

```
def generate_model_report(y_actual, y_predicted):
    print("Accuracy = ", accuracy_score(y_actual, y_predicted))
    print("Precision = ", precision_score(y_actual, y_predicted))
    print("Recall = ", recall_score(y_actual, y_predicted))
    print("F1 Score = ", f1_score(y_actual, y_predicted))
    pass
```

In [32]:

```
def generate_auc_roc_curve(clf, X_test):
    y_pred_proba = clf.predict_proba(X_test)[:, 1]
    fpr, tpr, thresholds = roc_curve(Y_test, y_pred_proba)
    auc = roc_auc_score(Y_test, y_pred_proba)
    plt.plot(fpr, tpr, label="AUC ROC Curve with Area Under the curve =" + str(auc))
    plt.legend(loc=4)
    plt.show()
    pass
```

In [33]:

```
df = pd.read_csv('data.csv')
```

In [34]:

```
df.tail()
```

Out[34]:

|        | Time     | V1        | V2        | V3       | V4       | V5       | V6       | V7       | V8       | V9 ...       | V21      | V22      |
|--------|----------|-----------|-----------|----------|----------|----------|----------|----------|----------|--------------|----------|----------|
| 284802 | 172786.0 | 11.881118 | 10.071785 | 9.834783 | 2.066656 | 5.364473 | 2.606837 | 4.918215 | 7.305334 | 1.914428 ... | 0.213454 | 0.111864 |

|        |          |           |           |          |          |          |          |          |          |          |     |          |          |
|--------|----------|-----------|-----------|----------|----------|----------|----------|----------|----------|----------|-----|----------|----------|
| 284803 | 172778.0 | -0.732799 | -0.055002 | 2.035009 | 0.738589 | 0.868229 | 1.058406 | 0.024309 | 0.294809 | 0.584809 | ... | 0.214925 | 0.924924 |
| 284804 | 172788.0 | 1.919565  | -0.301254 | 3.249640 | 0.557828 | 2.630515 | 3.031260 | 0.296827 | 0.708417 | 0.432454 | ... | 0.232045 | 0.578229 |
| 284805 | 172788.0 | -0.240440 | 0.530483  | 0.702510 | 0.689799 | 0.377961 | 0.623708 | 0.686180 | 0.679145 | 0.392087 | ... | 0.265245 | 0.800049 |
| 284806 | 172792.0 | -0.533413 | -0.189733 | 0.703337 | 0.506271 | 0.012546 | 0.649617 | 1.577006 | 0.414650 | 0.486180 | ... | 0.261057 | 0.643078 |

5 rows × 31 columns

In [35]:

```
X = df.loc[:, df.columns!='Class']
```

In [36]:

```
Y = df.loc[:, df.columns=='Class']
```


In [37]:

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.33, random_state=42)
```

In [38]:

```
ax = sns.countplot(x=df['Class'], data=df)
print(df['Class'].value_counts())
```

```
0    284315
1      492
Name: Class, dtype: int64
```



## Observation :

- Imbalanced Dataset

In [39]:

```
print(X_train.shape, Y_train.shape)
print(X_test.shape, Y_test.shape)
```


```
(190820, 30) (190820, 1)
(93987, 30) (93987, 1)
```

## Let's check imbalance in the train labels too

In [40]:

```
ax = sns.countplot(x=Y_train['Class'], data=Y_train)
print(Y_train['Class'].value_counts())
```

```
0    190477
1      343
Name: Class, dtype: int64
```



## Let's train a model

In [41]:

```
clf = LogisticRegression().fit(X_train, Y_train)
```

In [42]:

```
Y_Test_Pred = clf.predict(X_test)
```

In [43]:

```
pd.crosstab(Y_Test_Pred, Y_test['Class'], rownames=['Predicted'], colnames=['Actual'])
```

Out[43]:

| Actual | Predicted |     |
|--------|-----------|-----|
|        | 0         | 1   |
| 0      | 93792     | 47  |
| 1      | 46        | 102 |


In [44]:

```
generate_model_report(Y_test, Y_Test_Pred)
```

```
Accuracy = 0.9990105014523285
Precision = 0.6891891891891891
Recall = 0.6845637583892618
F1 Score = 0.6868686868686869
```

In [22]:

```
generate_auc_roc_curve(clf, X_test)
```



**Results are not bad at all.**

## Performing some tuning of class weights

In [45]:

```
weights = np.linspace(0.05, 0.95, 20)
gsc = GridSearchCV(
    estimator=LogisticRegression(),
    param_grid={
        'class_weight': [{0: x, 1: 1.0-x} for x in weights]
    },
    scoring='f1',
    cv=5
)

grid_result = gsc.fit(X_train, Y_train)
print("Best parameters : %s" % grid_result.best_params_)
```


Best parameters : {'class\_weight': {0: 0.4289473684210526, 1: 0.5710526315789475}}


In [48]:

```
data_out = pd.DataFrame({'score': grid_result.cv_results_['mean_test_score'], 'weight': weights })
data_out.plot(x='weight')
```

Out[48]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x119bcab50>





In [49]:

```
clf = LogisticRegression(**grid_result.best_params_).fit(X_train, Y_train)
Y_Test_Pred = clf.predict(X_test)
```

In [51]:

```
pd.crosstab(Y_Test_Pred, Y_test['Class'], rownames=['Predicted'], colnames=['Actual'])
```

Out[51]:

|           |   | Actual |    |
|-----------|---|--------|----|
|           |   | 0      | 1  |
| Predicted | 0 | 93802  | 59 |
|           | 1 | 36     | 90 |


In [52]:

```
generate_model_report(Y_test, Y_Test_Pred)
```

```
Accuracy = 0.9989892219136689
Precision = 0.7142857142857143
Recall = 0.6040268456375839
F1 Score = 0.6545454545454547
```

In [53]:

```
generate_auc_roc_curve(clf, X_test)
```



## Let's try using `class_weight='balanced'`

- $w_j = n/k \cdot n_j$ , where  $w_j$  is weight to class  $j$ ,  $n$  is the number of observations,  $n_j$  is the number of observations in class  $j$ ,  $k$  is the total number of classes.

In [54]:

```
clf = LogisticRegression(class_weight='balanced').fit(X_train, Y_train)
Y_Test_Pred = clf.predict(X_test)
```

In [55]:

```
pd.crosstab(Y_Test_Pred, Y_test['Class'], rownames=['Predicted'], colnames=['Actual'])
```

Out[55]:

| Actual    | 0     | 1   |
|-----------|-------|-----|
| Predicted |       |     |
| 0         | 90726 | 12  |
| 1         | 3112  | 137 |

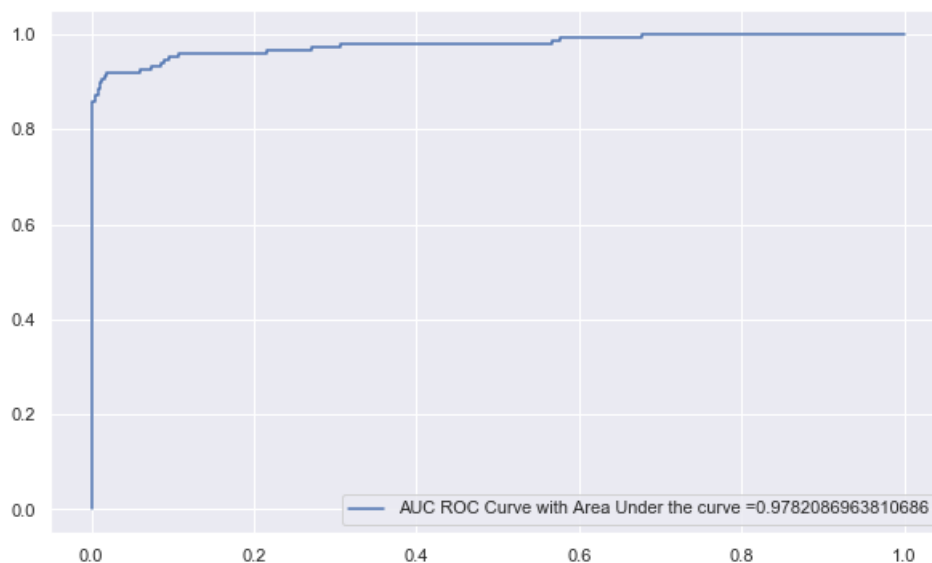
In [56]:

```
generate_model_report(Y_test, Y_Test_Pred)
```

```
Accuracy = 0.9667613606137019
Precision = 0.04216682056017236
Recall = 0.9194630872483222
F1 Score = 0.08063566804002353
```

In [57]:

```
generate_auc_roc_curve(clf, X_test)
```




**F1 score has dropped drastically since precision has dropped**

**Let's try the SMOTE technique**

In [58]:

```
ax = sns.countplot(x=Y_train['Class'], data=Y_train)
print(Y_train['Class'].value_counts())
```

```
0    190477
1       343
Name: Class, dtype: int64
```




In [60]:

```
sm = SMOTE(random_state=12, sampling_strategy='auto')
x_train_res, y_train_res = sm.fit_sample(X_train, Y_train)
```

In [61]:

```
ax = sns.countplot(x=y_train_res['Class'], data=y_train_res)
print(y_train_res['Class'].value_counts())
```

```
1    190477
0    190477
Name: Class, dtype: int64
```



In [62]:

```
clf = LogisticRegression().fit(x_train_res, y_train_res)
```

In [63]:

```
Y_Test_Pred = clf.predict(X_test)
```

In [65]:

```
pd.crosstab(Y_Test_Pred, Y_test['Class'], rownames=['Predicted'], colnames=['Actual'])
```

Out[65]:

| Actual    | 0     | 1   |
|-----------|-------|-----|
| Predicted |       |     |
| 0         | 92202 | 15  |
| 1         | 1636  | 134 |


In [66]:

```
generate_model_report(Y_test, Y_Test_Pred)
```

```
Accuracy = 0.9824337408364987
Precision = 0.07570621468926554
Recall = 0.8993288590604027
F1 Score = 0.13965607087024492
```

In [67]:

```
generate_auc_roc_curve(clf, X_test)
```



**The F1 Score is very low.**

In [97]:

```
pipe = make_pipeline(
    SMOTE(),
    LogisticRegression()
)

weights = np.linspace(0.005, 0.25, 10)

gsc = GridSearchCV(
    estimator=pipe,
    param_grid={
        'smote__sampling_strategy': weights
    },
    scoring='f1',
    cv=3
)
grid_result = gsc.fit(X_train, Y_train)

print("Best parameters : %s" % grid_result.best_params_)
weight_f1_score_df = pd.DataFrame({'score': grid_result.cv_results_['mean_test_score'],
                                     'smote__sampling_strategy': grid_result.cv_results_['smote__sampling_strategy']})
```

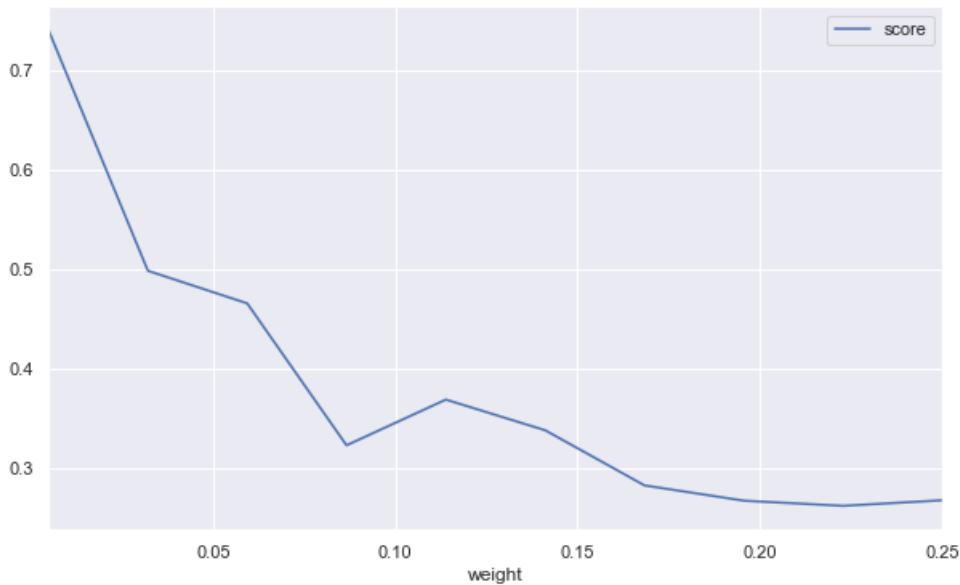


```
weight_f1_score_df.plot(x='weight',  
                        'weight': weights })
```

Best parameters : {'smote\_\_sampling\_strategy': 0.005}

Out[97]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1a3bd44890>



In [98]:

```
pipe = make_pipeline(  
    SMOTE(sampling_strategy=0.005),  
    LogisticRegression()  
)  
  
pipe.fit(X_train, Y_train)  
  
Y_Test_Pred = pipe.predict(X_test)
```

In [99]:

```
pd.crosstab(Y_Test_Pred, Y_test['Class'],  
            rownames=['Predicted'], colnames=['Actual'])
```

Out[99]:

| Predicted | Actual |     |
|-----------|--------|-----|
|           | 0      | 1   |
| 0         | 93761  | 37  |
| 1         | 77     | 112 |


In [100]:

```
generate_model_report(Y_test, Y_Test_Pred)
```

```
Accuracy = 0.9987870662964027  
Precision = 0.5925925925925926  
Recall = 0.7516778523489933  
F1 Score = 0.6627218934911243
```

In [96]:

```
generate_auc_roc_curve(clf, X_test)
```



## UNDERSAMPLING

In [102]:

```
minority_class_len = len(df[df['Class'] == 1])  
print(minority_class_len)
```

492

In [104]:

```
majority_class_len = len(df[df['Class'] == 0])  
print(majority_class_len)
```

284315

In [105]:

```
majority_class_indices = df[df['Class'] == 0].index
```

In [106]:

```
## Undersampling  
random_majority_indices = np.random.choice(majority_class_indices, minority_class_len, replace=False)  
print(len(random_majority_indices))
```

492

In [108]:

```
minority_class_indices = df[df['Class'] == 1].index
```

In [109]:

```
under_sample_indices = np.concatenate([minority_class_indices, random_majority_indices])
```

In [110]:


```
under_sample = df.loc[under_sample_indices]
```

In [111]:

```
sns.countplot(x='Class', data=under_sample)
```

Out[111]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1a390b2f50>



In [112]:

```
X = under_sample.loc[:, df.columns!='Class']  
Y = under_sample.loc[:, df.columns=='Class']  
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.33, random_state=42)  
clf = LogisticRegression().fit(X_train, Y_train)  
Y_Test_Pred = clf.predict(X_test)
```

In [113]:

```
generate_model_report(Y_test, Y_Test_Pred)
```

```
Accuracy = 0.8984615384615384  
Precision = 0.9337748344370861  
Recall = 0.8597560975609756  
F1 Score = 0.8952380952380953
```

In [114]:

```
generate_auc_roc_curve(clf, X_test)
```

